

Section 1 explains workplace expectations for the support and engineering teams. Remote staff must document important troubleshooting steps in a shared knowledge base. Escalations should include symptom description, affected document type, and observed retrieval behavior. Managers review repeated ingestion failures every Friday to identify training or tooling gaps. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Section 2 explains workplace expectations for the support and engineering teams. Remote staff must document important troubleshooting steps in a shared knowledge base. Escalations should include symptom description, affected document type, and observed retrieval behavior. Managers review repeated ingestion failures every Friday to identify training or tooling gaps. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Section 3 explains workplace expectations for the support and engineering teams. Remote staff must document important troubleshooting steps in a shared knowledge base. Escalations should include symptom description, affected document type, and observed retrieval behavior. Managers review repeated ingestion failures every Friday to identify training or tooling gaps. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Section 4 explains workplace expectations for the support and engineering teams. Remote staff must document important troubleshooting steps in a shared knowledge base. Escalations should include symptom description, affected document type, and observed retrieval behavior. Managers review repeated ingestion failures every Friday to identify training or tooling gaps. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Section 5 explains workplace expectations for the support and engineering teams. Remote staff must document important troubleshooting steps in a shared knowledge base. Escalations should include symptom description, affected document type, and observed retrieval behavior. Managers review repeated ingestion failures every Friday to identify training or tooling gaps. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Section 6 explains workplace expectations for the support and engineering teams. Remote staff must document important troubleshooting steps in a shared knowledge base. Escalations should include symptom description, affected document type, and observed retrieval behavior. Managers review repeated ingestion failures every Friday to identify training or tooling gaps. The system should preserve citation quality, page references, and chunk traceability

so a user can verify which source passage supported the answer.

Section 7 explains workplace expectations for the support and engineering teams. Remote staff must document important troubleshooting steps in a shared knowledge base. Escalations should include symptom description, affected document type, and observed retrieval behavior. Managers review repeated ingestion failures every Friday to identify training or tooling gaps. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Section 8 explains workplace expectations for the support and engineering teams. Remote staff must document important troubleshooting steps in a shared knowledge base. Escalations should include symptom description, affected document type, and observed retrieval behavior. Managers review repeated ingestion failures every Friday to identify training or tooling gaps. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Section 9 explains workplace expectations for the support and engineering teams. Remote staff must document important troubleshooting steps in a shared knowledge base. Escalations should include symptom description, affected document type, and observed retrieval behavior. Managers review repeated ingestion failures every Friday to identify training or tooling gaps. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Section 10 explains workplace expectations for the support and engineering teams. Remote staff must document important troubleshooting steps in a shared knowledge base. Escalations should include symptom description, affected document type, and observed retrieval behavior. Managers review repeated ingestion failures every Friday to identify training or tooling gaps. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Section 11 explains workplace expectations for the support and engineering teams. Remote staff must document important troubleshooting steps in a shared knowledge base. Escalations should include symptom description, affected document type, and observed retrieval behavior. Managers review repeated ingestion failures every Friday to identify training or tooling gaps. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Section 12 explains workplace expectations for the support and engineering teams. Remote staff must document important troubleshooting steps in a shared knowledge base. Escalations should include symptom description, affected document type, and observed retrieval behavior. Managers review repeated ingestion failures every Friday to identify training or tooling gaps. The system should preserve citation quality, page references, and chunk traceability

Data Handling Policy

Policy note 1 covers handling uploaded customer files in local environments. Agents must not commit uploaded files, model caches, or local vector data to source control. Test environments may store temporary PDFs and TXT samples only in approved project folders. When answers lack context, staff should state uncertainty instead of fabricating a response. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Policy note 2 covers handling uploaded customer files in local environments. Agents must not commit uploaded files, model caches, or local vector data to source control. Test environments may store temporary PDFs and TXT samples only in approved project folders. When answers lack context, staff should state uncertainty instead of fabricating a response. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Policy note 3 covers handling uploaded customer files in local environments. Agents must not commit uploaded files, model caches, or local vector data to source control. Test environments may store temporary PDFs and TXT samples only in approved project folders. When answers lack context, staff should state uncertainty instead of fabricating a response. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Policy note 4 covers handling uploaded customer files in local environments. Agents must not commit uploaded files, model caches, or local vector data to source control. Test environments may store temporary PDFs and TXT samples only in approved project folders. When answers lack context, staff should state uncertainty instead of fabricating a response. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Policy note 5 covers handling uploaded customer files in local environments. Agents must not commit uploaded files, model caches, or local vector data to source control. Test environments may store temporary PDFs and TXT samples only in approved project folders. When answers lack context, staff should state uncertainty instead of fabricating a response. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Policy note 6 covers handling uploaded customer files in local environments. Agents must not commit uploaded files, model caches, or local vector data to source control. Test environments may store temporary PDFs and TXT samples only in approved project folders. When

answers lack context, staff should state uncertainty instead of fabricating a response. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Policy note 7 covers handling uploaded customer files in local environments. Agents must not commit uploaded files, model caches, or local vector data to source control. Test environments may store temporary PDFs and TXT samples only in approved project folders. When answers lack context, staff should state uncertainty instead of fabricating a response. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Policy note 8 covers handling uploaded customer files in local environments. Agents must not commit uploaded files, model caches, or local vector data to source control. Test environments may store temporary PDFs and TXT samples only in approved project folders. When answers lack context, staff should state uncertainty instead of fabricating a response. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Policy note 9 covers handling uploaded customer files in local environments. Agents must not commit uploaded files, model caches, or local vector data to source control. Test environments may store temporary PDFs and TXT samples only in approved project folders. When answers lack context, staff should state uncertainty instead of fabricating a response. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Policy note 10 covers handling uploaded customer files in local environments. Agents must not commit uploaded files, model caches, or local vector data to source control. Test environments may store temporary PDFs and TXT samples only in approved project folders. When answers lack context, staff should state uncertainty instead of fabricating a response. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Operational Reference

Operational entry 1 summarizes recurring support situations. Common questions involve supported file types, the Qdrant storage path, and the Ollama endpoint. The default local vector database path is `backend/data/qdrant` and the default Ollama URL is `http://localhost:11434`. These facts are useful for verifying whether long-form retrieval remains grounded after pagination. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Operational entry 2 summarizes recurring support situations. Common questions involve supported file types, the Qdrant storage path, and the Ollama endpoint. The default local vector database path is backend/data/qdrant and the default Ollama URL is <http://localhost:11434>. These facts are useful for verifying whether long-form retrieval remains grounded after pagination. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Operational entry 3 summarizes recurring support situations. Common questions involve supported file types, the Qdrant storage path, and the Ollama endpoint. The default local vector database path is backend/data/qdrant and the default Ollama URL is <http://localhost:11434>. These facts are useful for verifying whether long-form retrieval remains grounded after pagination. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Operational entry 4 summarizes recurring support situations. Common questions involve supported file types, the Qdrant storage path, and the Ollama endpoint. The default local vector database path is backend/data/qdrant and the default Ollama URL is <http://localhost:11434>. These facts are useful for verifying whether long-form retrieval remains grounded after pagination. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Operational entry 5 summarizes recurring support situations. Common questions involve supported file types, the Qdrant storage path, and the Ollama endpoint. The default local vector database path is backend/data/qdrant and the default Ollama URL is <http://localhost:11434>. These facts are useful for verifying whether long-form retrieval remains grounded after pagination. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Operational entry 6 summarizes recurring support situations. Common questions involve supported file types, the Qdrant storage path, and the Ollama endpoint. The default local vector database path is backend/data/qdrant and the default Ollama URL is <http://localhost:11434>. These facts are useful for verifying whether long-form retrieval remains grounded after pagination. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Operational entry 7 summarizes recurring support situations. Common questions involve supported file types, the Qdrant storage path, and the Ollama endpoint. The default local

vector database path is backend/data/qdrant and the default Ollama URL is <http://localhost:11434>. These facts are useful for verifying whether long-form retrieval remains grounded after pagination. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Operational entry 8 summarizes recurring support situations. Common questions involve supported file types, the Qdrant storage path, and the Ollama endpoint. The default local vector database path is backend/data/qdrant and the default Ollama URL is <http://localhost:11434>. These facts are useful for verifying whether long-form retrieval remains grounded after pagination. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Operational entry 9 summarizes recurring support situations. Common questions involve supported file types, the Qdrant storage path, and the Ollama endpoint. The default local vector database path is backend/data/qdrant and the default Ollama URL is <http://localhost:11434>. These facts are useful for verifying whether long-form retrieval remains grounded after pagination. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.

Operational entry 10 summarizes recurring support situations. Common questions involve supported file types, the Qdrant storage path, and the Ollama endpoint. The default local vector database path is backend/data/qdrant and the default Ollama URL is <http://localhost:11434>. These facts are useful for verifying whether long-form retrieval remains grounded after pagination. The system should preserve citation quality, page references, and chunk traceability so a user can verify which source passage supported the answer.